

# Exercise Sheet 2, 2022

## 6.0 VU Advanced Database Systems

Nicolas Bschor (12132344)

8. April 2022

**Aufgabe 1 (Costs of MapReduce)** [1 Punkte] For this exercise we have to state a view consideration for the base scenario:

We define  $|R|$  as the number of input and  $\text{size}(|R|)$  as the I/O of reading the input.

- Communication Cost  $c$ :  $c = \text{size}(|R|)$  (Output very small  $\rightarrow$  not significantly affecting the Communication Cost)
  - Replication Rate  $rr$ :  $1 \gg rr \geq \frac{k}{|R|}$  (Replication Rate roughly around the ideal replication rate of  $\frac{k}{|R|}$  and much smaller than 1. Every Node emits a bit more than it's top  $k$  values, dependent on the data)
  - Reducer Size  $rs$ :  $1 \gg rs \geq k \cdot \text{nodes}$  (Every node emits a bit more than it's top  $k$  tuples to the key "None" is mapped on  $\geq k \cdot \text{nodes}$  elements)
- (a)
- Communication Cost  $c$ :  $c = \text{size}(|R|)$  (Only emit more tuples, so there is not more I/O)  $\rightarrow$  Same as base scenario
  - Replication Rate  $rr$ :  $rr = 1$  (For every input emit one tuple)  $\rightarrow$  Much higher than base scenario
  - Reducer Size  $rs$ :  $|R|$  (Every Node emits it's input on "None" key)
- (b) Every node emits  $\approx k$  values on a specific pre-reducer. After the pre-reducing the result is then mapped to the "None" key for the final reducing. We consider a useful distribution of the pre-reducers (e.g nodes 1-10 to reducer 1, 11-20 to reducer 2...).
- Communication Cost  $c$ :  $c = \text{size}(|R|) + [\text{prereducers} \cdot *k + k]$  (prereducers  $\cdot *k$ : I/O to write the result of pre-reducers (not high);  $k$ : write final result)  $\rightarrow$  A bit higher than base scenario.
  - Replication Rate  $rr$ : For both jobs, the same replication rate as in the base scenario.
  - Reducer Size  $rs$ :  $1 \gg rs \geq k \cdot \frac{\text{nodes}}{\text{pre-reducers}}$ ,  $1 \gg rs \geq k \cdot \text{pre-reducers}$ ,  $\rightarrow$  smaller than base scenario.

- Skew: Since we split the emitted keys evenly on the pre-reducers we have as much keys as pre-reducers and all should contain  $\approx$  the same size of the values  $\rightarrow$  No skew
- (c) • Communication Cost  $c$ :  $c = \text{size}(|R|) + \text{size}(k)$  (Now the number of  $k$  is high, so the I/O for writing can be a factor as well)
- Replication Rate  $rr$ :  $rr = \frac{k}{|R|}$  (We emit a lot more tuples since we search for more values)  $\rightarrow$  Much higher than base scenario
- Reducer Size  $rs$ :  $\text{nodes} * k \leq rr$  (since  $k$  is very high we have a lot of values mapped to "None")  $\rightarrow$  higher than base scenario

### Aufgabe 2 (Relational Operations) [2 Punkte]

(a)

$$F_1(R, S) := \sigma_{A+B=C}(R) \bowtie (\sigma_{B>3}(S) \ltimes \sigma_{C>B}(R)) = \sigma_{A+B=C \wedge C>B}(R) \bowtie \sigma_{B>3}(S)$$

Solution with a single map/ reducer pair.

#### Mapper:

When tuple  $t$  of  $S$ : If  $t.B > 3$  then emit the key/value pair:

$\langle (t_S.B, t_S.C), ("S", t_S) \rangle$

When tuple  $t$  of  $R$ : If  $t.C > t.B \wedge t.A + t.B = t.C$  then emit the key-value pair:

$\langle (t_R.B, t_R.C), ("R", t_R) \rangle$

#### Reducer:

Build two buffers of the values: Buffer one with the  $("R", t_R)$  values and buffer 2 with the  $("S", t_S)$  values. Then make a nested loop with the two buffers and emit the "cross product" of the two buffers. Key-value pair:  $(\_, (t_R \circ t_S))$

(b) 1) Job:

#### Mapper:

If tuple  $t \in T$ : emit  $\langle (t.B, t.C), t.A \rangle$

If tuple  $t \in S$ : emit  $\langle \text{NULL}, (t.B, t.C) \rangle$

#### Reducer:

If key = null: Sort values and emit  $\langle (t_1.B, t_1.C, t_2.B, t_2.C, \dots, t_n.B, t_n.C), \text{True} \rangle$  with  $t_1, t_2, \dots, t_n$  as the values.

If key  $\in A$ : Sort values and emit all lengths of the keys with the following form:

$\langle ((t_1.B, t_1.C)), \text{key} \rangle$

$\langle ((t_1.B, t_1.C), (t_2.B, t_2.C)), \text{key} \rangle$

...

$\langle ((t_1.B, t_1.C), (t_2.B, t_2.C), \dots, (t_n.B, t_n.C)), \text{key} \rangle$

2) Job:

**Reducer:**

If True in the Values: emit all other values

**Example:**  $U := [2,2], [1,1], [1,2]$

$T := [A1, 1,1], [A1, 1,2], [A1, 2,2], [A1, 3,3], [A2, 1,1], [A3, 3,3]$

1. Job Mapper:

From U:

Null  $\rightarrow [[2,2], [1,1], [1,2]]$

From T:  $A1 \rightarrow [[2,2], [1,1], [1,2]]$ ;  $A2 \rightarrow [1,1]$ ;  $A3 \rightarrow [3,3]$

1. Job Reducer:

From Null:  $([1,1], [1,2], [2,2]) \rightarrow \text{True}$

From A1:  $([1,1]) \rightarrow A1$ ;  $([1,1], [1,2]) \rightarrow \text{True}$ ;  $([1,1], [1,2], [2,2]) \rightarrow \text{True}$ ;

$([1,1], [1,2], [2,2], [3,3]) \rightarrow \text{True}$

From A2:  $[1,1] \rightarrow A2$

From A3:  $[3,3] \rightarrow A3$

2. Job Reducer:  $([1,1], [1,2], [2,2]) \rightarrow \text{True}$ ,  $A1 \Rightarrow \text{emit } A1$

**Aufgabe 3 (MapReduce in Practice)** [4 Punkte] Silence is gold

**Aufgabe 4 (Hive)** [4 Punkte] Silence is gold

**Aufgabe 5 (Spark in Scala)** [4 Punkte]

(a) File: "Exercise5\_SparkInScala.ipynb"

(b) First the csv file of "customers" is scanned. Then a filter process is executed and the filtered tuples are then exchanged and sorted afterwards. All these dependencies are **narrow dependencies**.

Now the csv file of "articles" is scanned. Then it is filtered (**narrow transformation**) and broadcasted afterwards (**narrow dependency**).

Then the csv file of "transactions" is scanned. It's also filtered through a (**narrow transformation**) and afterwards BroadcastHashJoined with the results from "articles" (**wide dependency**).

The joined result of "articles" and "transactions" is then exchanged and sorted (**narrow dependency**).

Now a **narrow dependent** SortMergeJoin can be applied the joined tuples of "articles" & "transactions" and "customers".

With the join result two **narrow dependent** HachAggregate can be executed, to emit the final result.